

project

June 28, 2020

0.1 EDA

```
[25]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn.linear_model import LinearRegression
#1
df = pd.read_csv('kc_house_data.csv')

# Correlation Plot Heatmap
plt.figure(figsize= (12, 12))
sns.heatmap(df.corr())
df.corr(method='pearson')
```

```
[25]:
```


	id	price	bedrooms	bathrooms	sqft_living	sqft_lot	\
id	1.000000	-0.016772	0.001150	0.005162	-0.012241	-0.131911	
price	-0.016772	1.000000	0.308787	0.525906	0.701917	0.089876	
bedrooms	0.001150	0.308787	1.000000	0.514508	0.578212	0.032471	
bathrooms	0.005162	0.525906	0.514508	1.000000	0.755758	0.088373	
sqft_living	-0.012241	0.701917	0.578212	0.755758	1.000000	0.173453	
sqft_lot	-0.131911	0.089876	0.032471	0.088373	0.173453	1.000000	
floors	0.018608	0.256804	0.177944	0.502582	0.353953	-0.004814	
waterfront	-0.002727	0.266398	-0.006834	0.063744	0.103854	0.021632	
view	0.011536	0.397370	0.080008	0.188386	0.284709	0.074900	
condition	-0.023803	0.036056	0.026496	-0.126479	-0.059445	-0.008830	
grade	0.008188	0.667951	0.356563	0.665838	0.762779	0.114731	
sqft_above	-0.010799	0.605368	0.479386	0.686668	0.876448	0.184139	
sqft_basement	-0.005193	0.323799	0.302808	0.283440	0.435130	0.015418	
yr_built	0.021617	0.053953	0.155670	0.507173	0.318152	0.052946	
yr_renovated	-0.016925	0.126424	0.018389	0.050544	0.055308	0.007686	
zipcode	-0.008211	-0.053402	-0.154092	-0.204786	-0.199802	-0.129586	
lat	-0.001798	0.306692	-0.009951	0.024280	0.052155	-0.085514	
long	0.020672	0.022036	0.132054	0.224903	0.241214	0.230227	

sqft_living15	-0.002701	0.585241	0.393406	0.569884	0.756402	0.144763
sqft_lot15	-0.138557	0.082845	0.030690	0.088303	0.184342	0.718204

	floors	waterfront	view	condition	grade	\
id	0.018608	-0.002727	0.011536	-0.023803	0.008188	
price	0.256804	0.266398	0.397370	0.036056	0.667951	
bedrooms	0.177944	-0.006834	0.080008	0.026496	0.356563	
bathrooms	0.502582	0.063744	0.188386	-0.126479	0.665838	
sqft_living	0.353953	0.103854	0.284709	-0.059445	0.762779	
sqft_lot	-0.004814	0.021632	0.074900	-0.008830	0.114731	
floors	1.000000	0.023755	0.028814	-0.264075	0.458794	
waterfront	0.023755	1.000000	0.401971	0.016611	0.082888	
view	0.028814	0.401971	1.000000	0.045999	0.251728	
condition	-0.264075	0.016611	0.045999	1.000000	-0.146896	
grade	0.458794	0.082888	0.251728	-0.146896	1.000000	
sqft_above	0.523989	0.072109	0.167609	-0.158904	0.756073	
sqft_basement	-0.245715	0.080559	0.277078	0.173849	0.168220	
yr_built	0.489193	-0.026153	-0.053636	-0.361592	0.447865	
yr_renovated	0.006427	0.092873	0.103951	-0.060788	0.014261	
zipcode	-0.059541	0.030272	0.084622	0.002888	-0.185771	
lat	0.049239	-0.014306	0.005871	-0.015102	0.113575	
long	0.125943	-0.041904	-0.078107	-0.105877	0.200341	
sqft_living15	0.280102	0.086507	0.280681	-0.093072	0.713867	
sqft_lot15	-0.010722	0.030781	0.072904	-0.003126	0.120981	

	sqft_above	sqft_basement	yr_built	yr_renovated	zipcode	\
id	-0.010799	-0.005193	0.021617	-0.016925	-0.008211	
price	0.605368	0.323799	0.053953	0.126424	-0.053402	
bedrooms	0.479386	0.302808	0.155670	0.018389	-0.154092	
bathrooms	0.686668	0.283440	0.507173	0.050544	-0.204786	
sqft_living	0.876448	0.435130	0.318152	0.055308	-0.199802	
sqft_lot	0.184139	0.015418	0.052946	0.007686	-0.129586	
floors	0.523989	-0.245715	0.489193	0.006427	-0.059541	
waterfront	0.072109	0.080559	-0.026153	0.092873	0.030272	
view	0.167609	0.277078	-0.053636	0.103951	0.084622	
condition	-0.158904	0.173849	-0.361592	-0.060788	0.002888	
grade	0.756073	0.168220	0.447865	0.014261	-0.185771	
sqft_above	1.000000	-0.052156	0.424037	0.023251	-0.261570	
sqft_basement	-0.052156	1.000000	-0.133064	0.071233	0.074725	
yr_built	0.424037	-0.133064	1.000000	-0.224907	-0.347210	
yr_renovated	0.023251	0.071233	-0.224907	1.000000	0.064325	
zipcode	-0.261570	0.074725	-0.347210	0.064325	1.000000	
lat	-0.001199	0.110414	-0.148370	0.029350	0.266742	
long	0.344842	-0.144546	0.409993	-0.068321	-0.564259	
sqft_living15	0.731767	0.200443	0.326377	-0.002695	-0.279299	
sqft_lot15	0.195077	0.017550	0.070777	0.007944	-0.147294	

	lat	long	sqft_living15	sqft_lot15
id	-0.001798	0.020672	-0.002701	-0.138557
price	0.306692	0.022036	0.585241	0.082845
bedrooms	-0.009951	0.132054	0.393406	0.030690
bathrooms	0.024280	0.224903	0.569884	0.088303
sqft_living	0.052155	0.241214	0.756402	0.184342
sqft_lot	-0.085514	0.230227	0.144763	0.718204
floors	0.049239	0.125943	0.280102	-0.010722
waterfront	-0.014306	-0.041904	0.086507	0.030781
view	0.005871	-0.078107	0.280681	0.072904
condition	-0.015102	-0.105877	-0.093072	-0.003126
grade	0.113575	0.200341	0.713867	0.120981
sqft_above	-0.001199	0.344842	0.731767	0.195077
sqft_basement	0.110414	-0.144546	0.200443	0.017550
yr_built	-0.148370	0.409993	0.326377	0.070777
yr_renovated	0.029350	-0.068321	-0.002695	0.007944
zipcode	0.266742	-0.564259	-0.279299	-0.147294
lat	1.000000	-0.135371	0.048679	-0.086139
long	-0.135371	1.000000	0.335626	0.255586
sqft_living15	0.048679	0.335626	1.000000	0.183515
sqft_lot15	-0.086139	0.255586	0.183515	1.000000




Most Positive: Sqft, Bedrooms, Bathrooms Most Negative: Zipcode, Lat, Long

(2). Sale numbers Vs. (years,months) and Sale prices correlation Vs. (years,months)


```
[26]: df['month'] = pd.DatetimeIndex(df['date']).month
df['year'] = pd.DatetimeIndex(df['date']).year
month = df['month'].value_counts()
year = df['year'].value_counts()
sns.barplot(year.index.tolist(),year)
plt.title("2014 and 2015 sale numbers")
```

[26]: Text(0.5, 1.0, '2014 and 2015 sale numbers')




```
[27]: sns.barplot(month.index.tolist(),month)  
plt.title("sale numbers in each month")
```

```
[27]: Text(0.5, 1.0, 'sale numbers in each month')
```



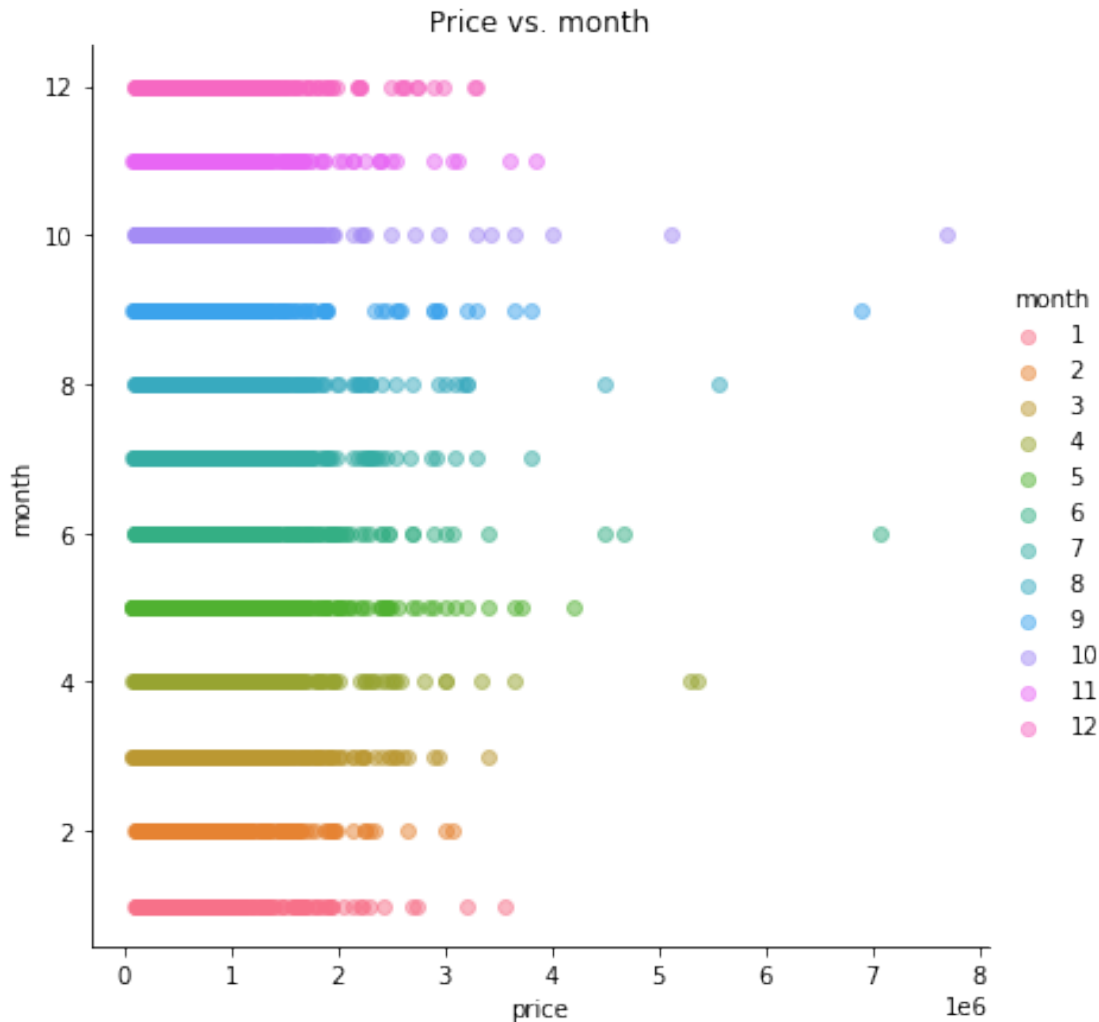
```
[28]: g=sns.FacetGrid(df,hue='year',height=6)
g.map(plt.scatter,'price','year',alpha=0.5)
g.add_legend()
plt.title("Price vs. Year")
```

```
[28]: Text(0.5, 1.0, 'Price vs. Year')
```



```
[29]: g=sns.FacetGrid(df,hue='month',height=6)
g.map(plt.scatter,'price','month',alpha=0.5)
g.add_legend()
plt.title("Price vs. month")
```

```
[29]: Text(0.5, 1.0, 'Price vs. month')
```



```
[30]: print("Price correlation with year: ",df['price'].corr(df['year']))
      print("Price correlation with month: ",df['price'].corr(df['month']))
```

Price correlation with year: 0.003727139624315499

Price correlation with month: -0.009928289245273971

```
[31]: #3 from sklearn.linear_model import LinearRegression
fig, (ax, box, sq, yr) = plt.subplots(4, figsize=(10,10))
plt.subplots_adjust(hspace = .5)
# Price
ax = sns.violinplot(ax = ax, x = df['price'])
print(np.percentile(df['price'], [25, 50, 75]))
ax.set(xlabel = 'Price')

# 'bedrooms', 'bathrooms', 'floors', 'condition'
```

```

box = df.boxplot(ax = box, column = ['bedrooms', 'bathrooms', 'floors', 'condition', 'grade'])
box.set_ylim([0,15])

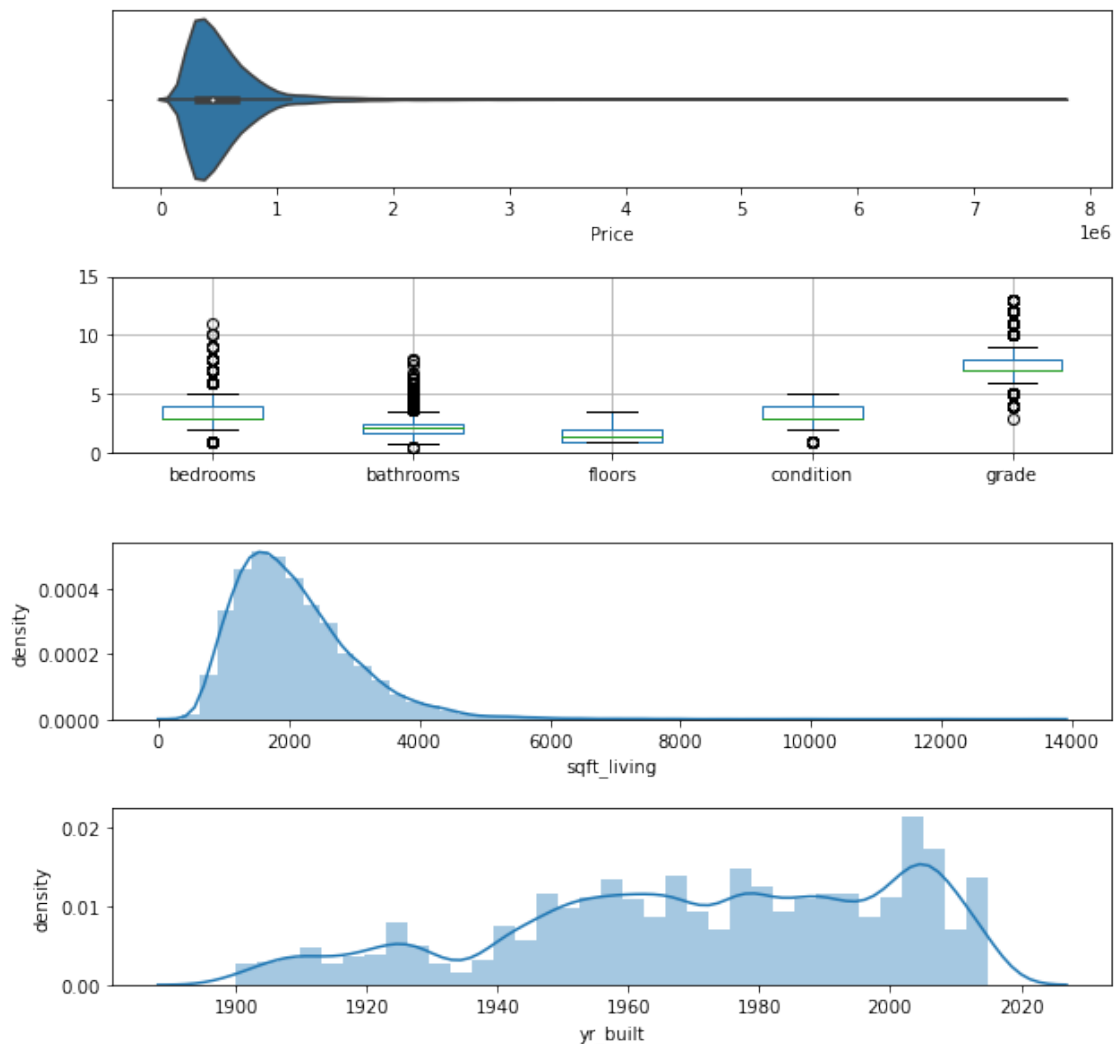
# Square feet living room
sq = sns.distplot(df.sqft_living, ax = sq)
sq.set(ylabel = 'density')

# Year built
yr = sns.distplot(df.yr_built, ax = yr)
yr.set(ylabel = 'density')

```

[322000. 450000. 645000.]

[31]: [Text(0, 0.5, 'density')]



1 The first graph shows the distribution of prices in a violin plot. We can tell the 25-75% quartile is between \$322,000 and \$645,000 2 The second plot shows the box plots of bedrooms, bathrooms, floors, condition, grade. The medians are: Bedrooms ~ 3 Bathrooms ~ 2.5 Floors ~ 2 Condition ~ 3 Grade ~ 7 3 The third plot shows the distribution of square foot in living room. This plot is skewed with the most being ~1800 sqft 4 The last plot is the distribution of houses built over time. There has been a recent phase of construction in the 2000s, which means many houses are newly built and in decent condition.

(4). Create the scoring function for 'Grade' with accuracy:70%


```
[32]: X=df[['price','bedrooms','bathrooms','sqft_living15','sqft_lot15','floors','waterfront',"condition",
        "yr_renovated"]]
y=df['grade']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
```

```
[33]: clf = RandomForestClassifier(n_estimators=100)
clf.fit(X_train,y_train)
y_predict=clf.predict(X_test)
```

```
[34]: variables = pd.Series(clf.feature_importances_,index=X.columns).
        ↪sort_values(ascending=False)
variables
```

```
[34]: price                0.210859
sqft_living15            0.195388
yr_built                 0.144138
sqft_lot15               0.130732
zipcode                 0.097867
bathrooms                0.096445
floors                   0.037984
bedrooms                 0.037329
condition                0.022929
view                     0.015756
yr_renovated             0.009387
waterfront               0.001187
dtype: float64
```

```
[35]: ax=sns.barplot(x=variables.index, y=variables,palette=sns.color_palette("RdBu",12),
        ↪12))
ax.set_xticklabels(ax.get_xticklabels(), rotation=40, ha="right")
# Add labels to your graph
plt.xlabel('Variables Importance Score')
plt.ylabel('weight')
plt.title("Visualizing Variables Weight")
plt.show()
```




```
[36]: print("Scoring function accuracy:", metrics.accuracy_score(y_test, y_predict))
```

Scoring function accuracy: 0.7037037037037037

```
[37]: #5
import seaborn as sns

wealthy = df.loc[df['price'] >= 3000000]

plt.title("Wealthy Areas in King County by Million $")
ax = sns.scatterplot(x=wealthy.lat, y=wealthy.long, hue=wealthy.price)
```




0.2 Modeling

0.2.1 Linear Regression

```
[38]: X=df[['bedrooms','bathrooms','sqft_living15','grade','sqft_lot15','floors','waterfront',"condi
      "yr_renovated"]]
      y=df['price']
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
```

```
[39]: plt.figure(figsize=(15,10))
      plt.tight_layout()
      sns.distplot(df['price'])
      plt.title("distribution of house prices")
```

```
[39]: Text(0.5, 1.0, 'distribution of house prices')
```



```
[40]: reg = LinearRegression()
      reg.fit(X_train,y_train)
      coeff_df = pd.DataFrame(reg.coef_, X.columns, columns=['Coefficient'])
      coeff_df
```

```
[40]:
```

	Coefficient
bedrooms	-8017.909375
bathrooms	110699.450542
sqft_living15	89.997854
grade	159993.490592
sqft_lot15	-0.237560
floors	18715.992830
waterfront	582458.379984
condition	24022.402672
yr_built	-4049.375666
zipcode	21.457573
view	48132.877243
yr_renovated	10.387863

```
[41]: y_predit = reg.predict(X_test)
      accurate_rate=1-np.mean(np.abs(y_predit-y_test)/y_test)
      print("Accuracy: ",accurate_rate)
```

Accuracy: 0.6995068223297689


```
[42]: result = pd.DataFrame({'Actual price': y_test, 'Predicted price': y_predit})  
      result.head(8)
```

```
[42]:
```

	Actual price	Predicted price
16644	570000.0	948422.541466
11196	427800.0	410592.227532
534	204000.0	106154.319955
16261	270000.0	503446.838729
17002	760369.0	672881.313048
625	289500.0	185084.386431
10450	330000.0	287006.767480
13494	570000.0	384449.655879

```
[43]: g = sns.FacetGrid(result,height=8)  
      g.map(plt.scatter,'Actual price','Predicted price',alpha=0.5)  
      plt.plot([0,2500000],[0,2500000],ls='--',color='orange',label='accuracy line')  
      plt.title("Actual price Vs. Predicted price")  
      plt.legend()  
      print("Model accuracy: ",accurate_rate)
```

Model accuracy: 0.6995068223297689



Linear regression is a model to find possible W , in “ $Y = XW + \text{error}$ ” which has minimum Mean squared error (MSE). This linear regression model accuracy rate is around 66%.


0.2.2 Random Forest Model

```
[44]: X=df[['bedrooms','bathrooms','sqft_living15','grade','sqft_lot15','floors','waterfront','condi
      "yr_renovated"]]
      y=df['price']
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
      clf = RandomForestRegressor(n_estimators=100)
      clf.fit(X_train,y_train)
      y_predict=clf.predict(X_test)
```

```
[45]: variables = pd.Series(clf.feature_importances_,index=X.columns).
      ↪sort_values(ascending=False)
      variables
```


```
[45]: grade                0.505252
      yr_built             0.109656
      zipcode              0.097298
      sqft_living15        0.072033
      bathrooms            0.069986
      sqft_lot15           0.047110
      waterfront           0.041807
      view                 0.019622
      bedrooms             0.013280
      condition            0.008312
      yr_renovated         0.007956
      floors               0.007688
      dtype: float64
```

```
[46]: ax=sns.barplot(x=variables, y=variables.index,palette=sns.color_palette("PRGn",10),
      ↪10))
      ax.figure.set_size_inches(12,10)
      # Add labels to your graph
      plt.xlabel('weight ')
      plt.ylabel('scoring variables')
      plt.title("Visualizing Variables Weight")
      plt.show()
```



```
[47]: g = sns.FacetGrid(result,height=8)
g.map(plt.scatter,'Actual price','Predicted price',alpha=0.5)
plt.plot([0,2500000],[0,2500000],ls='--',color='orange',label='accuracy line')
plt.title("Actual price Vs. Predicted price")
plt.legend()
print("Model accuracy: ",accurate_rate)
```

Model accuracy: 0.6995068223297689



```
[48]: accurate_rate=1-np.mean(np.abs(y_predit-y_test)/y_test)
      print("Random Forest accuracy:",accurate_rate)
```

Random Forest accuracy: 0.8274844176466372

Random forest regression is to select random samples and build decision trees for each sample. Then, Perform a vote for each predicted result and select the prediction result with the most votes as the final prediction. The Random forest model has accuracy rate around 81%.

0.2.3 K Nearest Neighbors

```
[49]: from sklearn import neighbors
X=df[['bedrooms','bathrooms','floors','grade','sqft_living15','waterfront',"condition","yr_built",
      "yr_renovated"]]
y=df['price']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)

model = neighbors.KNeighborsRegressor(n_neighbors=10)
model.fit(X_train, y_train)
preds = model.predict(X_test)
```


```
[50]: result = pd.DataFrame({'Actual price': y_test, 'Predicted price': preds})
result.head(8)
```

```
[50]:
```

	Actual price	Predicted price
11743	250000.0	286345.0
3810	532000.0	422565.0
9408	375000.0	435894.0
10552	195000.0	203434.5
5363	1600000.0	657495.0
8963	394999.0	407325.0
18741	999000.0	433090.0
7915	305000.0	298410.0

```
[51]: g = sns.FacetGrid(result,height=8)
g.map(plt.scatter,'Actual price','Predicted price',alpha=0.5)
plt.plot([0,2500000],[0,2500000],ls='--',color='orange',label='accuracy line')
plt.title("Actual price Vs. Predicted price")
plt.legend()
```

```
[51]: <matplotlib.legend.Legend at 0x19dc50ec748>
```



```
[52]: print('Nearest Neighbors Accuracy: ', model.score(X_test, y_test))
```

Nearest Neighbors Accuracy: 0.4211697320021205

KNN works by finding the distances between a query and all the examples in the data, selecting the specified number examples (K) closest to the query, then votes for the most frequent label (in the case of classification) or averages the labels (in the case of regression).

This method is not the best because it yielded a ~45% accuracy.